

# Week 2 Notes

## Lectures 3 & 4

University of Massachusetts, Amherst, CS389

### 1 Neural Networks (lecture 3)

**Definition 1.1** (Multi-dimensional data). A *multi-dimensional data* has  $n$  number of data and  $d$  number of features and can be written as<sup>1</sup>:

$$X = [X^{(1)}, X^{(2)}, \dots, X^{(n)}]^T \quad (1.1)$$

where  $X^{(i)}$  is the  $i^{\text{th}}$  data<sup>2</sup>. A single data,  $X^{(i)}$  is a  $d$ -dimensional vector which has  $d$  number of features. For example, the features for a weather data may include humidity, temperature, air speed, etc. Every data is represented as a row and each feature as a column. Formally we write a single data as:

$$X^{(i)} = [X_1^{(i)}, X_2^{(i)}, \dots, X_d^{(i)}] \quad (1.2)$$

Putting all this together, our multi dimensional data is a matrix with dimensions  $(n \times d)$ :

$$\begin{bmatrix} X_1^{(1)} & \dots & X_d^{(1)} \\ \dots & \dots & \dots \\ X_1^{(n)} & \dots & X_d^{(n)} \end{bmatrix} \quad (1.3)$$

#### 1.1 Perceptron

**Definition 1.2** (Perceptron).




Figure 1: A single-layer perceptron

where the vertices  $X_i$  are the inputs, the edges  $W_i$  are the weighted connections, the vertex  $b$  is the bias term, and the vertex  $Y'$  is the output. We can write this mathematically as:

$$Y' = \phi(W^T X + b) \quad (1.4)$$

where  $\phi$  is a non-linearity function, also known as an activation function (topic of next week).

<sup>1</sup>Note that this array is transposed because  $X$  is a row vector and each  $X^{(i)}$  is a  $d$ -dimensional column vector

<sup>2</sup>The reason why we use the  $X^{(i)}$  notation is because  $X^{(i)}$  is a vector and sometimes we need to refer to the elements of that vector [2]

## 1.2 Classification

In a classification problem, we want to learn to predict discrete classes which the input belongs to.

**Definition 1.3** (Binary classification). Is a supervised learning algorithm that *categorizes the data into one of two classes*.

Recall that a linear model is defined as:

$$\hat{y} = \phi \left( \sum_{j=1}^m w_j \cdot X_j + b \right) = \phi(W^T X + b) \quad (1.5)$$

where the output  $\hat{y} \in \mathbb{R}$  and we can apply a boolean function (i.e. the *sign* activation function) so that the output  $\hat{y} \in \{-1, +1\}$ . Given a fixed model, the binary classifier is defined as:

$$\hat{y} = \text{sign}(W^T X + b) = \begin{cases} +1 & \text{if } W^T X + b > 0 \\ -1 & \text{otherwise} \end{cases} \quad (1.6)$$

As can be seen in Figure 2, the points that are on the same side as the normal vector to the hyperplane is the positive half-spaces are classified as positive. Likewise, the other half-space is negative and all points are classified as negative [2].




Figure 2: Binary Classification Example. The  $\theta$  is the same variable as  $W$  [2]

## 1.3 Linear Separability

**Theorem 1.** If the dataset,  $D$ , is linearly separable, then the perceptron algorithm is guaranteed to find a linear separator [2]

How would one *formally describe linear separability*? This is beyond the scope for this class, but you can refer to this [MIT lecture notes](#) for a detailed explanation. The intuition is if the shortest distance of a point to the hyperplane (the norm) is positive for all points then the dataset is classified correctly. You can use Figure 2 to help with this visualization.

## 1.4 Regression

**Definition 1.4** (Regression). Predict continuous outputs ( $\hat{y} \in \mathbb{R}$ ) that are close to the true values

## 2 Stochastic Gradient Descent implementation (lecture 4)

Hopefully, we now have a mathematical understanding and an intuition of all the main components of a supervised machine learning model. We can now start implementing a very simple SGD. Recall that in a SGD, we want to update the parameters  $W$  and  $b$  after every single training data.

---

**Algorithm 1** Stochastic Gradient Descent algorithm

---

```
1: for  $e$  in range Epoch do                                     ▷ number of epochs
2:   for  $X^{(i)}$  in  $X$  do                                           ▷ loop through entire dataset
3:      $\hat{y}^{(i)} = F_W(X^{(i)}) = X^{(i)}W^T + b$                      ▷ our model's prediction for input  $X^{(i)}$ 
4:      $\text{Loss}^{(i)} = \text{MSE}(y^{(i)}, \hat{y}^{(i)}) = \frac{1}{2} \sum (y^{(i)} - \hat{y}^{(i)})^2$  ▷ calculate loss of prediction with actual
5:      $W = W - \alpha \nabla_W \text{Loss}^{(i)}$                                ▷ update weight
6:      $b = b - \alpha \nabla_b \text{Loss}^{(i)}$                                ▷ update bias
7:   end for
8: end for
```

---

### 2.1 Calculating Gradient of Loss

We define the gradient of loss as  $\nabla_W \text{Loss}$ , which can also be written as  $\frac{\partial L}{\partial W}$ , such that:

$$\nabla_W \text{Loss} = \frac{\partial L}{\partial W} = \left( \frac{\partial L}{\partial W_1}, \frac{\partial L}{\partial W_2}, \dots, \frac{\partial L}{\partial W_m} \right) \quad (2.1)$$

Using the chain rule we have:

$$\frac{\partial L}{\partial W_i} = \frac{\partial \hat{y}}{\partial W_i} \frac{\partial L}{\partial \hat{y}} \quad (2.2)$$

Now recall that  $\hat{y} = W^T X + b$  and  $L = \frac{1}{2} \sum (y^{(i)} - \hat{y}^{(i)})^2$ , therefore:

$$\frac{d\hat{y}}{dW_1} = X_1 \quad (2.3)$$

$$\frac{d\hat{y}}{db} = 1 \quad (2.4)$$

## References

- [1] Cooper.
- [2] MIT Open Learning Library, 6.036, Spring 2020, [https://openlearninglibrary.mit.edu/assets/courseware/v1/2481f8f2964716032b134db99e369b81/asset-v1:MITx+6.036+1T2019+type@asset+block/notes\\_chapter\\_Introduction.pdf](https://openlearninglibrary.mit.edu/assets/courseware/v1/2481f8f2964716032b134db99e369b81/asset-v1:MITx+6.036+1T2019+type@asset+block/notes_chapter_Introduction.pdf)